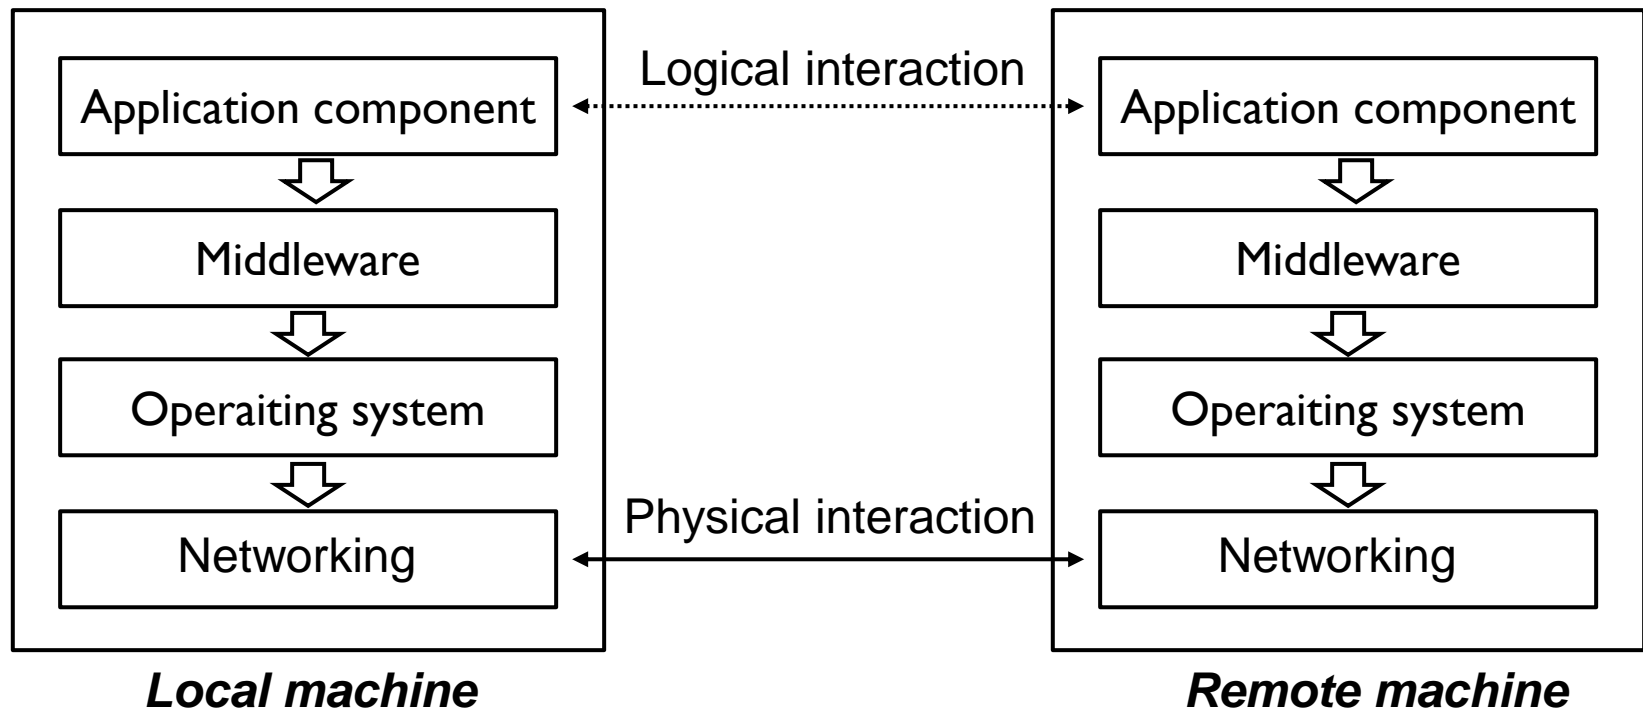


Development of Large Systems (RPC: remote procedure call)

Andrea Corradini
aco@cphnusinnes.dk

Remote procedure call (RPC)

- ▶ RPC supports client-server computing
 - ▶ servers offer a set of operations through a service interface
 - ▶ clients call these operations as if they were local




RPC features

- ▶ Allows a process executing on a machine to invoke a procedure in a process executing on another machine on a network
- ▶ Is a form of synchronous inter-process communication
 - ▶ requesting program is suspended until the result of the remote procedure is calculated and returned
- ▶ Based on request-response protocol (i.e. client-server)
 - ▶ initiated by the client sending a request message to a remote server to execute a given procedure with supplied parameters
- ▶ Abstracts procedure calls between processes and simplifies the process of writing distributed applications
 - ▶ preserving the syntax of a local procedure call while transparently initiating network communication
- ▶ Requires an IDL to describe the procedures that can be called over the network as a standard vocabulary for exchanging information

RPC flow

- Programs using RPC are compiled into executables that include a stub acting as representative of the remote procedure code



RPC semantics in case of failures

▶ Failure types

- ▶ the client is unable to locate the server
- ▶ the request message from the client to the server is lost
- ▶ the reply message from the server to the client is lost
- ▶ the server crashes after receiving a request
- ▶ the client crashes after sending a request

Failure 1: client cannot locate the server

- ▶ We distinguish between five different classes of failures that can occur in RPC systems
 - ▶ the client is unable to locate the server
 - ▶ the request message from the client to the server is lost
 - ▶ the reply message from the server to the client is lost
 - ▶ the server crashes after receiving a request
 - ▶ the client crashes after sending a request
- ▶ Common solutions: report failure, signal handlers, exception handling, etc.

Failure 2: lost request messages

- ▶ We distinguish between five different classes of failures that can occur in RPC systems
 - ▶ the client is unable to locate the server
 - ▶ the request message from the client to the server is lost
 - ▶ the reply message from the server to the client is lost
 - ▶ the server crashes after receiving a request
 - ▶ the client crashes after sending a request

Failure 3: lost reply messages

- ▶ We distinguish between five different classes of failures that can occur in RPC systems
 - ▶ the client is unable to locate the server
 - ▶ the request message from the client to the server is lost
 - ▶ the reply message from the server to the client is lost
 - ▶ the server crashes after receiving a request
 - ▶ the client crashes after sending a request

Failure 2 and 3: treatment

- ▶ Client waits for reply message, resends request upon timeout
 - ▶ client cannot tell whether the request or reply was lost
- ▶ Client can safely resend a request for idempotent operations
 - ▶ idempotent operations produce the same result when executed repeatedly
- ▶ For non idempotent operations, client can add sequence numbers to requests for the server to tell a retransmitted request from the original one
 - ▶ server need keep track of the most recently received sequence number from each client
 - ▶ server will not carry out a retransmitted request, but sends a reply to the client

Failure 4: server crash after receiving request

- ▶ We distinguish between five different classes of failures that can occur in RPC systems
 - ▶ the client is unable to locate the server
 - ▶ the request message from the client to the server is lost
 - ▶ the reply message from the server to the client is lost
 - ▶ the server crashes after receiving a request
 - ▶ the client crashes after sending a request

Failure 4: treatments

- ▶ The client cannot tell if the crash occurred before or after the request is carried out
 - ▶ request arrives, but the server crashes before sending the reply
 - client must retransmit the request
 - ▶ request arrives and is carried out, but the server crashes before sending the reply
 - server somehow has to report back to the client e.g. by raising an exception

Failure 4: approaches

- ▶ Give up immediately and report back failure (*at most once semantics*)
 - ▶ RPC is carried out at most one time, but possibly none at all
- ▶ Keep trying until a reply has been received then give it to the client (*at least once semantics*)
 - ▶ RPC is carried out at least one time, or possibly more
- ▶ Do not guarantee anything
 - ▶ when a server crashes, the client gets no help/promises
 - ▶ the RPC may have been carried out any number of times
- ▶ Ideal is *exactly once semantics* but usually no way to arrange it

Failure 5: client crashes

- ▶ We distinguish between five different classes of failures that can occur in RPC systems
 - ▶ the client is unable to locate the server
 - ▶ the request message from the client to the server is lost
 - ▶ the reply message from the server to the client is lost
 - ▶ the server crashes after receiving a request
 - ▶ the client crashes after sending a request

Failure 5: dealing with orphans

- ▶ **Extermination:** client keeps log entry before sending an RPC and kills off the orphan when it comes back up
 - ▶ expensive for memory storage
 - ▶ orphans may make RPC i.e., grand orphans are created
- ▶ **Reincarnation:** when a client reboots, it broadcasts a new epoch number; when server receives the broadcast, it kills all the computations running on behalf of the client
- ▶ **Gentle reincarnation:** variant of reincarnation, when computations killed only for clients that cannot be located
- ▶ **Expiration:** each RPC is given an amount of time T to do job
 - ▶ if it cannot finish, it must explicitly ask for another quantum T but right amount of time T is unknown

RPC protocols

- ▶ Connection-oriented protocol or connectionless protocol?
- ▶ Standard general-purpose protocol or one specifically designed for RPC?

XML-RPC

- ▶ Protocol uses a subset of XML vocabulary as encoding mechanism and interface definition language
 - ▶ client specifies a procedure name and parameters in the XML request
 - ▶ server returns a response or a failure in the XML response format
- ▶ Protocol uses the HTTP protocol to pass information from client to a server
 - ▶ XML messages are the payload of the HTTP response and request

XML-RPC data: basic types

- ▶ XML-RPC specification defines six basic data types
 - ▶ used in passing parameters, return values, and error messages
 - ▶ can be combined to create two more compound data types to support almost any data types in any programming language
- ▶ Basic data type:
 - ▶ 32-bit integer like e.g. `<int>23097</int>`
 - ▶ 64-bit double like e.g. `<double>-23.097</double>`
 - ▶ boolean like e.g. `<boolean>1</boolean>`
 - ▶ string like e.g. `<string>Hello World!</string>`
 - ▶ date and time like e.g. `<dateTime.iso8601>20210321T13:08:29</dateTime.iso8601>`
 - ▶ binary values like e.g. `<base64>GpyBpcyBhHNob3J0Ncm3k0luZy</base64>`

XML-RPC compound data: array

- ▶ Multidimensional arrays can also be created by defining arrays inside `<value>` `</value>` tags

```
<value>
  <array>
    <data>
      <value><boolean>1</boolean></value>
      <value><string>It is true</string></value>
      <value><int>23097</int></value>
    </data>
  </array>
</value>
```

XML-RPC compound data: struct

- ▶ Structs are not objects

```
<value>
  <struct>
    <member>
      <name>givenName</name>
      <value><string>Andrea</string></value>
    </member>
    <member>
      <name>familyName</name>
      <value><string>Corradini</string></value>
    </member>
  </struct>
</value>
```

Correspondence with types in Java

XML-RPC Type	Type	Java Type
<i4> or <int>	Four byte signed integer	int
<boolean>	0 (false) or 1 (true)	boolean
<string>	string	String
<double>	double-precision signed floating point number	double
<dateTime.iso8601>	date/time	N/A
<base64>	base64-encoded binary	byte[] or String
<array>	array of elements	array object
<struct>	record of elements	object

Coding example: client request

- ▶ Method call encoded as message sent from client to server look like:

```
<?xml version="1.0" encoding="ISO-8859-1"?>
<methodCall>
  <methodName> sumhandler.sum</methodName>
  <params>
    <param> <value><int>23</int></value> </param>
    <param> <value><int>9</int></value> </param>
  </params>
</methodCall>
```

Coding example I: server response

- ▶ Messages returned from server to client look like:

```
<?xml version="1.0" encoding="ISO-8859-1"?>  
<methodResponse>  
  <params>  
    <param> <value><int>32</int></value> </param>  
  </params>  
</methodResponse>
```

Coding example II: server error response

```
<?xml version="1.0" encoding="ISO-8859-1"?>
<methodResponse>
  <fault>
    <value>
      <struct>
        <member>
          <name>fault</name>
          <value><int>99</int></value>
        </member>
        <member>
          <name>faultString</name>
          <value><string>No such method</string></value>
        </member>
      </struct>
    </value>
  </fault>
</methodResponse>
```

RPC limitations

- ▶ Limited number of data types
- ▶ No provision for passing objects as it is not compatible with objects
- ▶ No type checking of array values; mixed type not forbidden
- ▶ Strings allow only ASCII characters
- ▶ No check for duplicate names in struct
- ▶ No representation of NaN for double
- ▶ Limited security since firewalls are bypassed using HTTP

Coding exercise: XML-RPC

- ▶ Do the tutorial on XML-RPC at https://www.tutorialspoint.com/xml-rpc/xml_rpc_examples.htm
 - ▶ it is a remote call for a simple addition operation
 - ▶ add more operations
 - ▶ make the entire project more OOP (i.e., split responsibilities between objects)
 - ▶ display XML messages exchanged between client and server
 - ▶ display XML error messages exchanged between client and server
 - ▶ display the time it takes to run a method remotely and compare it with the time it take to run it locally
- ▶ (elective) Do your own client and server using Java sockets to send back and forth messages in HTTP

Headers for HTTP (for elective part)

- ▶ When method calls and method responses are sent using HTTP, messages require certain specific headers

Method call (from client)

POST / HTTP 1.1

Host: server host name

User-Agent: software making the request

Content-Type: text/xml

Content-Length: nr of bytes in payload

... here the payload ...

Method response (from server)

HTTP 200 OK

Content-Type: text/xml

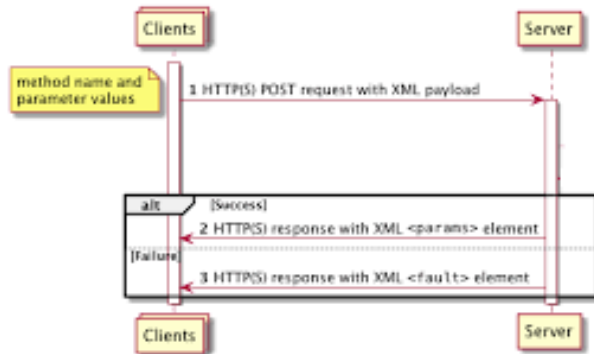
Content-Length: nr of bytes in payload

... here the payload...

RPC-XML: how to code it in Java (for elective part)

CLIENT

- 1 - Build an XML element *methodCall* that names the method to call and provides its parameters
- 2 - Send a POST request whose payload (content) is the XML element just built over a socket connection



- 7 - Receive the response message, parse the XML element returned, and use the outcome in the client code

SERVER

- 3 - Receive the request and use HTTP header Content-Length to read the XML payload
- 4 - Parse the XML element, extract method name, and retrieve the its parameters
- 5 - Search for the desired method and, if found, invoke it with the given parameters using reflection
- 6 - If method executes successfully, package the return value in *methodResponse* XML element (otherwise package a fault message), and send that message back to the client over the socket